

# Communication Paradigms in IoT

Understanding the fundamental communication paradigms is crucial for designing robust and efficient IoT systems. Each approach offers distinct advantages and trade-offs concerning latency, scalability, and resource consumption. This card introduces the three primary models and their fundamental characteristics.



## Publish/Subscribe

**Broker-based, exemplified by MQTT.** Decoupled in time, space, and synchronisation. Ideal for high-fanout telemetry and event distribution.

- Asynchronous communication
- High scalability for many receivers
- Decoupling simplifies system design



## Request/Response

**Client sends a request and awaits a synchronous reply.** Simpler, natural fit for CRUD operations (e.g., HTTP, CoAP).

- Synchronous blocking model
- Well-suited for transaction-based operations
- Easily understood with standard web practices



## Persistent Duplex

**Maintains a bidirectional channel (e.g., WebSockets).** Enables real-time, interactive communication flows.

- Full-duplex channel persistence
- Low latency once established
- Resource-intensive connection management

Design choices are heavily influenced by factors such as latency requirements, scalability needs, energy efficiency, and firewall traversal capabilities. Selecting the optimal paradigm dictates the success of an IoT deployment.

# MQTT Fundamentals: The IoT Messaging Standard

MQTT (Message Queuing Telemetry Transport) is a lightweight, publish-subscribe messaging protocol designed for constrained devices and low-bandwidth, high-latency, or unreliable networks. Its broker-centric architecture ensures efficient message routing and scalability.

## Core Concepts

- **Broker-centric:** All messages pass through a central broker, which manages client connections, subscriptions, and message routing.
- **Lightweight:** Minimal overhead over TCP/IP, critical for battery-powered devices and low-bandwidth networks.
- **Clients:** Devices act as publishers (sending data) or subscribers (receiving data).
- **Topics:** Hierarchical naming structure for message routing, supporting wildcards (+ for single level, # for multi-level).
- **Session Control:** Persistent sessions maintain subscriptions and queued messages; Clean sessions discard them upon disconnect.

## Quality of Service (QoS) Levels



### QoS 0: At Most Once

Best-effort delivery. Messages are sent once, with no acknowledgment of receipt. Fastest, lowest overhead. **No guarantee of delivery.**

### QoS 1: At Least Once

Messages are guaranteed to arrive, but duplicates are possible. Requires acknowledgment from receiver. **Guaranteed delivery, potential duplicates.**

### QoS 2: Exactly Once

Highest guarantee; messages arrive exactly once. Achieved via a four-step handshake between sender and receiver. **Guaranteed delivery without duplication.**

Additional features include **Retained Messages** (broker stores the last message on a topic for new subscribers) and **Last Will & Testament** (a message published by the broker on behalf of a client if it disconnects unexpectedly), ensuring system resilience and state management.

# MQTT Scaling and Security Considerations

Implementing MQTT in production environments requires careful attention to scalability and robust security measures to handle vast numbers of connections and sensitive data.

## Transport & Scaling

- **Standard Ports:** TCP 1883 for unencrypted, TLS 8883 for encrypted communication.
- **MQTT over WebSocket:** Allows browser-based clients to connect, leveraging existing web infrastructure and simplifying firewall traversal.
- **Clustering:** Distributing broker workload across multiple instances to ensure high availability and load balancing.
- **Topic Sharding:** Assigning different topics to different brokers or broker clusters to manage high message volumes.
- **Session Replication:** Ensuring client sessions and state are maintained across multiple broker nodes for fault tolerance.

## Security Protocols

### TLS Encryption

Transport Layer Security (TLS) encrypts data in transit, protecting against eavesdropping and tampering. Essential for sensitive data over untrusted networks.

### Client Authentication

Verifies the identity of connecting clients using X.509 certificates or username/password credentials before granting access to the broker.

### Access Control Lists (ACLs)

Define granular permissions for clients, controlling precisely which topics they can publish to or subscribe from, preventing unauthorised data access.

## Trade-offs

- **Broker Centralisation:** Presents a single point of failure without proper clustering and load distribution.
- **Storage Load:** Retained messages can consume significant broker memory if not effectively managed and limited.
- **QoS2 Overhead:** The four-step handshake, while ensuring exactly-once delivery, introduces latency and increased network traffic compared to lower QoS levels.

# CoAP Core Concepts: The RESTful IoT Protocol

CoAP (Constrained Application Protocol) is a specialised web transfer protocol designed for resource-constrained devices and low-power, lossy networks (LLNs) in the Internet of Things. It brings RESTful principles to the edge.

## Key Characteristics

| → UDP-based  | → RESTful   | → Compact Header  |
|--|---|---|
| Utilises User Datagram Protocol, making it efficient due to minimal overhead but requiring application-level mechanisms to ensure reliability. | Adheres to Representational State Transfer principles, offering familiar methods (GET, POST, PUT, DELETE) for resource interaction. | Extremely small header sizes (typically <10 bytes), which is ideal for environments with stringent bandwidth constraints. |

## CoAP Message Types

- **Confirmable (CON):** Requires an acknowledgment (ACK) from the receiver, implementing application-layer reliability.
- **Non-confirmable (NON):** Does not require an acknowledgment, used for less critical data or periodic telemetry where occasional loss is acceptable.
- **Acknowledgment (ACK):** Confirms successful receipt of a CON message.
- **Reset (RST):** Indicates that a receiver has received a message but cannot process it.

CoAP uses **Tokens** for correlating requests and responses, and **Message IDs** for duplicate detection and rudimentary reliability management.

## Key Extensions for Advanced IoT Functionality

### CoAP Observe

Enables clients to subscribe to resource changes, receiving push notifications from the server without repetitive polling. This is crucial for event-driven architectures.

### Blockwise Transfer

Facilitates the transfer of larger payloads by breaking them into smaller blocks, overcoming the inherent constraints of limited UDP packet sizes in LLNs.

# CoAP Reliability and Security

Despite its UDP foundation, CoAP incorporates sophisticated mechanisms to ensure message reliability and implements stringent security protocols tailored for constrained environments.

## Reliability Mechanisms

- **Confirmable (CON) Retransmissions:** CON messages are retransmitted if no acknowledgment is received, employing an increasing exponential backoff timer to manage network congestion.
- **Congestion Control:** Achieved through duplicate suppression and randomised retransmission timers to prevent network overload and ensure fairness.
- **Duplicate Suppression:** Message IDs are used extensively to identify and discard duplicate messages, preventing unintended resource changes.

## Security Implementations



### DTLS (Datagram Transport Layer Security)

The UDP-equivalent of TLS, providing end-to-end security for CoAP messages by encrypting the communication channel and ensuring data integrity and confidentiality.



### OSCORE (Object Security for Constrained RESTful Environments)

Provides end-to-end object security by encrypting and authenticating the CoAP payload itself. This is vital as it protects the data even when traversing intermediate proxies or gateways.

## CoAP Proxies and Interoperability

Proxies are critical for bridging CoAP networks to the broader internet:

- **CoAP-HTTP Mapping:** Proxies can translate CoAP requests into HTTP and vice-versa, facilitating seamless cloud integration and access by standard web services.
- **Caching:** Improves performance and reduces network traffic by storing frequently accessed CoAP resources locally at the proxy.
- **Translation:** Enables effective communication between energy-efficient CoAP devices and high-bandwidth HTTP-based services.

# HTTP(S) in IoT: Ubiquitous Yet Challenging

HTTP and its secure counterpart, HTTPS, are universally supported protocols, making them attractive for certain IoT applications despite their inherent verbosity and overhead, which challenges resource-constrained devices.

## Advantages in IoT Context

### Ubiquity

Universally supported across networks, operating systems, and cloud platforms, simplifying integration and reducing development complexity.

### Firewall-friendly

Standard ports (80/443) are typically open in corporate and consumer networks, significantly easing network traversal and deployment.

### Easy REST APIs

The well-defined REST architectural style is widely understood and supported for interacting with web services and cloud backends.

## Major Drawbacks for Constrained Devices

- Verbose Headers:** Large header sizes introduce significant overhead, making the ratio of useful data to total transmitted data very low for small IoT payloads.
- TCP/TLS Handshake Overhead:** Initial connection establishment for each request is computationally and time-intensive, negatively impacting latency and device energy consumption.
- Request-Response Polling:** The synchronous model often requires devices to poll for updates, which is inefficient compared to asynchronous push models like MQTT or CoAP Observe.

## HTTP/2 Enhancements

HTTP/2 mitigates some traditional HTTP limitations, offering benefits for less constrained IoT gateways and applications:

- Multiplexing Streams:** Allows multiple requests and responses over a single persistent TCP connection, improving efficiency over individual transactions.
- Header Compression (HPACK):** Reduces overhead by compressing HTTP headers using a shared indexed table, significantly improving data efficiency.

HTTP(S) is best suited for device management, configuration updates, and cloud APIs where devices are not severely constrained and high interactivity is not the primary requirement.

# WebSockets: Real-Time Bidirectional Communication

WebSockets provide a persistent, full-duplex communication channel over a single TCP connection. This capability is ideal for delivering real-time, instantaneous data exchange in applications such as monitoring dashboards and interactive control systems.

## Key Characteristics

### HTTP Upgrade Mechanism

The connection is initiated as a standard HTTP request on port 80 or 443, which is then "upgraded" via a handshake to a persistent WebSocket connection. This facilitates easy traversal of firewalls and proxies.

### Persistent Duplex Channel

Once established, the connection maintains an open, bidirectional TCP channel for continuous data flow, eliminating the overhead of repeated TCP handshakes and request/response cycles.

## Features and Technical Details

- Lower Per-Message Overhead:** After the initial setup, the framing overhead for subsequent messages is minimal compared to the verbose headers of repeated HTTP requests.
- Text/Binary Frames:** Supports efficient transfer of both human-readable text and raw binary data using frame-based messaging.
- Ping/Pong Keepalive:** Built-in mechanism to ensure the connection remains active, measure latency, and detect network issues without needing application-level heartbeats.

## Primary Use Cases in IoT



### Real-time Dashboards

Displaying live sensor data, device status, or analytics with minimal delay to operational staff.



### Interactive Control Loops

Enabling immediate command and control of physical devices from a user interface, requiring rapid feedback.



### M2M Instant Messaging

Facilitating instantaneous, direct communication paths between machine-to-machine (M2M) interfaces.

While highly effective for real-time applications, WebSockets require a full TCP/IP stack, rendering them unsuitable for the most ultra-low-power microcontrollers.

# Protocol Selection Criteria for IoT

Choosing the correct messaging protocol is critical for overall system performance, cost efficiency, and reliability. The decision must be an objective process guided by application requirements and device constraints.

## Comparative Analysis of Core Protocols

| Overhead (per message) | Highest                | Lowest                | Low                 | Medium                   |
|------------------------|------------------------|-----------------------|---------------------|--------------------------|
| Transport              | TCP (request/response) | UDP (app reliability) | TCP (broker-based)  | TCP (duplex)             |
| Latency                | High (polling)         | Low (CON retrans.)    | Very Low (push)     | Very Low (duplex)        |
| Energy Efficiency      | Lowest                 | Highest               | High                | Medium                   |
| Firewall Traversal     | Excellent (HTTP ports) | Challenging (UDP)     | Moderate (TCP port) | Excellent (HTTP upgrade) |

## Heuristic Mapping for IoT Use Cases



### Telemetry & Eventing

MQTT for efficient, high-volume, and asynchronous data streams from numerous devices.



### Local Device Control

CoAP for constrained local networks and direct, resource-oriented device interaction with minimal overhead.



### Public APIs & Cloud Integration

HTTP(S) for its universality and ease of integration with external web services and traditional cloud interfaces.



### Real-time Dashboards & UI

WebSockets for instantaneous, persistent, and bidirectional updates to user-facing applications.

Engineers must balance the desire for low energy consumption (CoAP, MQTT) with the need for simplicity and firewall compatibility (HTTP, WebSockets).

# Hybrid Architectures and Bridging in IoT

Complex, heterogeneous IoT deployments necessitate hybrid architectures that strategically leverage the strengths of different protocols. This is achieved primarily through the deployment of intelligent gateways and bridging solutions at the network edge.

## Role of Gateways: The Protocol Translator

- **Protocol Translation:** Converting messages between disparate IoT protocols (e.g., CoAP payload into an MQTT message structure).
- **Data Aggregation:** Collecting and batching data from multiple edge devices to reduce the number of transmissions to the cloud.
- **Security Enforcement:** Implementing security policies, authenticating devices, and managing encryption at the edge before data egress.
- **Edge Computing:** Pre-processing data, performing real-time analytics, filtering noise, and enabling local decision-making to reduce latency.

## Common Bridging Patterns



### CoAP (Sensors) → MQTT Broker

Constrained edge devices use CoAP to communicate efficiently with a local gateway. The gateway then bridges this data to a centralised MQTT broker for wider distribution or ingestion by cloud services.



### MQTT → HTTP REST API

Data published to MQTT topics is consumed by a service running on the gateway or cloud. This service exposes the data via structured HTTP REST APIs, enabling web applications to interact with real-time IoT data using familiar methods.

## Challenges in Bridging Architectures

- **Semantic Mapping:** Translating RESTful resource URIs (CoAP) to hierarchical topics (MQTT) and vice versa requires a robust, well-defined mapping strategy.
- **QoS Mismatch:** Reconciling the different Quality of Service levels and reliability models between protocols (e.g., UDP-based CoAP reliability versus TCP-based MQTT reliability).
- **Retries & Idempotency:** Ensuring that retries due to network failures do not result in duplicate actions or state changes across protocol boundaries.

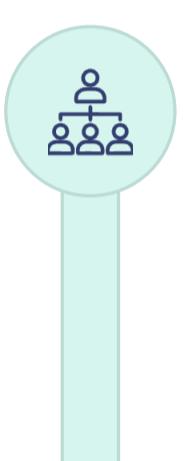
# Best Practices for IoT Messaging

Optimising IoT messaging requires a holistic approach that goes beyond mere protocol selection, focusing on efficient payload encoding, structured naming, robust reliability, stringent security, and comprehensive observability.



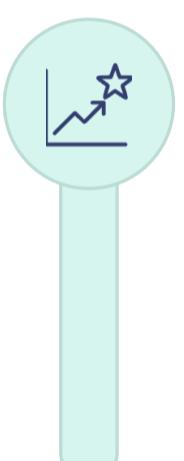
## Payload Encoding for Efficiency

Use compact binary formats like CBOR (Concise Binary Object Representation) or Protobuf (Protocol Buffers) for maximum efficiency on constrained networks. Reserve JSON for interoperability where bandwidth permits.



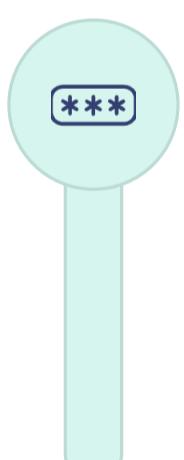
## Structured Naming Conventions

Adopt structured, hierarchical, and human-readable identifiers for topics (MQTT) or resources (CoAP). Consistency in naming significantly aids in management, debugging, and system scalability.



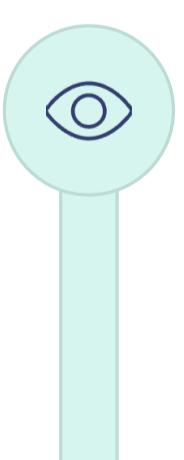
## Reliability Optimisation

Minimise QoS overhead (favour QoS 0 or 1 over QoS 2 if loss is tolerable). Implement robust exponential backoff strategies for retries, and ensure all critical operations are idempotent.



## Mandatory Security Protocols

Mandate TLS/DTLS for transport encryption. Consider OSCORE for end-to-end object security. Ensure unique device IDs and implement secure over-the-air (OTA) update mechanisms.



## Observability and System Health

Implement comprehensive monitoring for key metrics such as message latency, packet loss, message duplication, and broker throughput to diagnose issues and ensure continuous system health and operational integrity.

Adhering to academic principles, the protocol choice must always be empirically validated under real-world constraints, including power budget, latency requirements, reliability targets, and firewall policies. This ensures that the chosen solution is truly fit for purpose in the target deployment environment.